

## Solving problems by searching

From AIMA slides

1

## Outline

- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms

2

## Problem-solving agents

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
         state, some description of the current world state
         goal, a goal, initially null
         problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then do
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```


3

## Example: Romania

- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest
- **Formulate goal:**
  - be in Bucharest
- **Formulate problem:**
  - **states:** various cities
  - **actions:** drive between cities
- **Find solution:**
  - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

4

## Example: Romania



5


## Problem types

- **Deterministic, fully observable** → **single-state problem**
  - Agent knows exactly which state it will be in; solution is a sequence
- **Non-observable** → **sensorless problem (conformant problem)**
  - Agent may have no idea where it is; solution is a sequence
- **Nondeterministic and/or partially observable** → **contingency problem**
  - percepts provide **new** information about current state
  - often **interleave** search, execution
- **Unknown state space** → **exploration problem**

6

## Example: vacuum world


- Single-state, start in #5.  
Solution?



7

## Example: vacuum world


- Single-state, start in #5.  
Solution? [Right, Suck]



8

## Example: vacuum world

- Sensorless, start in {1,2,3,4,5,6,7,8} e.g., Right goes to {2,4,6,8}  
Solution? [Right, Suck, Left, Suck]




- Contingency
  - Nondeterministic: Suck may dirty a clean carpet
  - Partially observable: location, dirt at current location...
  - Percept: [L, Clean], i.e., start in #5 or #7
- Solution?

9

## Example: vacuum world

- Sensorless, start in {1,2,3,4,5,6,7,8} e.g., Right goes to {2,4,6,8}  
Solution? [Right, Suck, Left, Suck]



- Contingency
  - Nondeterministic: Suck may dirty a clean carpet
  - Partially observable: location, dirt at current location.
  - Percept: [L, Clean], i.e., start in #5 or #7
- Solution? [Right, if dirt then Suck]

10

## Single-state problem formulation

A problem is defined by four items:

- initial state e.g., "at Arad"
  - actions or successor function  $S(x)$  = set of action-state pairs
    - e.g.,  $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
  - goal test, can be
    - explicit, e.g.,  $x = \text{"at Bucharest"}$
    - implicit, e.g.,  $\text{Checkmate}(x)$
  - path cost (additive)
    - e.g., sum of distances, number of actions executed, etc.
    - $c(x, a, y)$  is the step cost, assumed to be  $\geq 0$
- A solution is a sequence of actions leading from the initial state to a goal state


11

## Selecting a state space

- Real world is absurdly complex
  - state space must be abstracted for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
  - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, any real state "in Arad" must get to some real state "in Zerind"
- (Abstract) solution =
  - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

12


## Vacuum world state space graph



- states?
- actions?
- goal test?
- path cost?

13


## Vacuum world state space graph



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

14


## Example: The 8-puzzle



- states?
- actions?
- goal test?
- path cost?

15

## Example: The 8-puzzle




- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

[Note: optimal solution of  $n$ -Puzzle family is NP-hard]

16

## Example: robotic assembly



- states?: real-valued coordinates of robot joint angles parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute

17


## Tree search algorithms

- Basic idea:
  - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. **expanding** states)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
```


18

## Tree search example




19

## Tree search example



20

## Tree search example



21

## Implementation: general tree search


```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    fringe ← INSERT ALL(EXPAND(node, problem), fringe)
```

```
function EXPAND(node, problem) returns a set of nodes
  successors ← the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s ← a new NODE
    PARENT-NODE[s] ← node; ACTION[s] ← action; STATE[s] ← result
    PATH-COST[s] ← PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s] ← DEPTH[node] + 1
    add s to successors
  return successors
```

22

## Implementation: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost**  $g(x)$ , **depth**



- The **Expand** function creates new nodes, filling in the various fields and using the **SUCCESSORFN** of the problem to create the corresponding states.

23

## Search strategies

- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
  - completeness**: does it always find a solution if one exists?
  - time complexity**: number of nodes generated
  - space complexity**: maximum number of nodes in memory
  - optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
  - $b$ : maximum branching factor of the search tree
  - $d$ : depth of the least-cost solution
  - $m$ : maximum depth of the state space (may be  $\infty$ )

24


## Uninformed search strategies

- **Uninformed** search strategies use only the information available in the problem definition
- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search

25

## Breadth-first search


- Expand shallowest unexpanded node
- **Implementation:**
  - *fringe* is a FIFO queue, i.e., new successors go at end



26

## Breadth-first search


- Expand shallowest unexpanded node
- **Implementation:**
  - *fringe* is a FIFO queue, i.e., new successors go at end



27

## Breadth-first search


- Expand shallowest unexpanded node
- **Implementation:**
  - *fringe* is a FIFO queue, i.e., new successors go at end



28

## Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
  - *fringe* is a FIFO queue, i.e., new successors go at end



29

## Properties of breadth-first search

- **Complete?** Yes (if  $b$  is finite)
- **Time?**  $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$
- **Space?**  $O(b^{d+1})$  (keeps every node in memory)
- **Optimal?** Yes (if cost = 1 per step)
- **Space** is the bigger problem (more than time)

30


## Uniform-cost search

- Expand least-cost unexpanded node
- Implementation:**
  - fringe* = queue ordered by path cost
- Equivalent to breadth-first if step costs all equal
- Complete?** Yes, if step cost  $\geq \epsilon$
- Time?** # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{\lceil \frac{C^*}{\epsilon} \rceil})$  where  $C^*$  is the cost of the optimal solution
- Space?** # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{\lceil \frac{C^*}{\epsilon} \rceil})$
- Optimal?** Yes – nodes expanded in increasing order of  $g(n)$

31

## Depth-first search


- Expand deepest unexpanded node
- Implementation:**
  - fringe* = LIFO queue, i.e., put successors at front



32

## Depth-first search


- Expand deepest unexpanded node
- Implementation:**
  - fringe* = LIFO queue, i.e., put successors at front



33

## Depth-first search


- Expand deepest unexpanded node
- Implementation:**
  - fringe* = LIFO queue, i.e., put successors at front



34

## Depth-first search

- Expand deepest unexpanded node
- Implementation:**
  - fringe* = LIFO queue, i.e., put successors at front



35

## Depth-first search


- Expand deepest unexpanded node
- Implementation:**
  - fringe* = LIFO queue, i.e., put successors at front



36

## Depth-first search


- Expand deepest unexpanded node
- Implementation:
  - fringe* = LIFO queue, i.e., put successors at front



37

## Depth-first search


- Expand deepest unexpanded node
- Implementation:
  - fringe* = LIFO queue, i.e., put successors at front



38

## Depth-first search


- Expand deepest unexpanded node
- Implementation:
  - fringe* = LIFO queue, i.e., put successors at front



39

## Depth-first search


- Expand deepest unexpanded node
- Implementation:
  - fringe* = LIFO queue, i.e., put successors at front



40

## Depth-first search


- Expand deepest unexpanded node
- Implementation:
  - fringe* = LIFO queue, i.e., put successors at front



41

## Depth-first search


- Expand deepest unexpanded node
- Implementation:
  - fringe* = LIFO queue, i.e., put successors at front



42

## Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - fringe = LIFO queue, i.e., put successors at front



43

## Properties of depth-first search

- Complete?** No: fails in infinite-depth spaces, spaces with loops
  - Modify to avoid repeated states along path
    - complete in finite spaces
- Time?**  $O(b^m)$ : terrible if  $m$  is much larger than  $d$ 
  - but if solutions are dense, may be much faster than breadth-first
- Space?**  $O(bm)$ , i.e., linear space!
- Optimal?** No

44

## Depth-limited search

= depth-first search with depth limit  $l$   
i.e., nodes at depth  $l$  have no successors

- Recursive implementation:

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
  RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)
function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
  cutoff-occurred? ← false
  if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
  else if DEPTH[node] = limit then return cutoff
  else for each successor in EXPAND(node, problem) do
    result ← RECURSIVE-DLS(successor, problem, limit)
    if result = cutoff then cutoff-occurred? ← true
    else if result ≠ failure then return result
  if cutoff-occurred? then return cutoff else return failure
```

45

## Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  inputs: problem, a problem
  for depth ← 0 to ∞ do
    result ← DEPTH-LIMITED-SEARCH(problem, depth)
    if result ≠ cutoff then return result
```


46

## Iterative deepening search $l = 0$



47


## Iterative deepening search $l = 1$



48




## Iterative deepening search / =2



49

## Iterative deepening search / =3



50

## Iterative deepening search

- Number of nodes generated in a depth-limited search to depth  $d$  with branching factor  $b$ :  

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$
- Number of nodes generated in an iterative deepening search to depth  $d$  with branching factor  $b$ :  

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$
- For  $b = 10$ ,  $d = 5$ :
  - $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
  - $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$
- Overhead =  $(123,456 - 111,111)/111,111 = 11\%$

51

## Properties of iterative deepening search

- Complete?** Yes
- Time?**  $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space?**  $O(bd)$
- Optimal?** Yes, if step cost = 1

52


## Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{(C^*/\epsilon)})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{(C^*/\epsilon)})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

53

## Repeated states

- Failure to detect repeated states can turn a linear problem into an exponential one!



54



## Graph search

```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERT-ALL(EXPAND(node, problem), fringe)
```

55



## Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms

56